


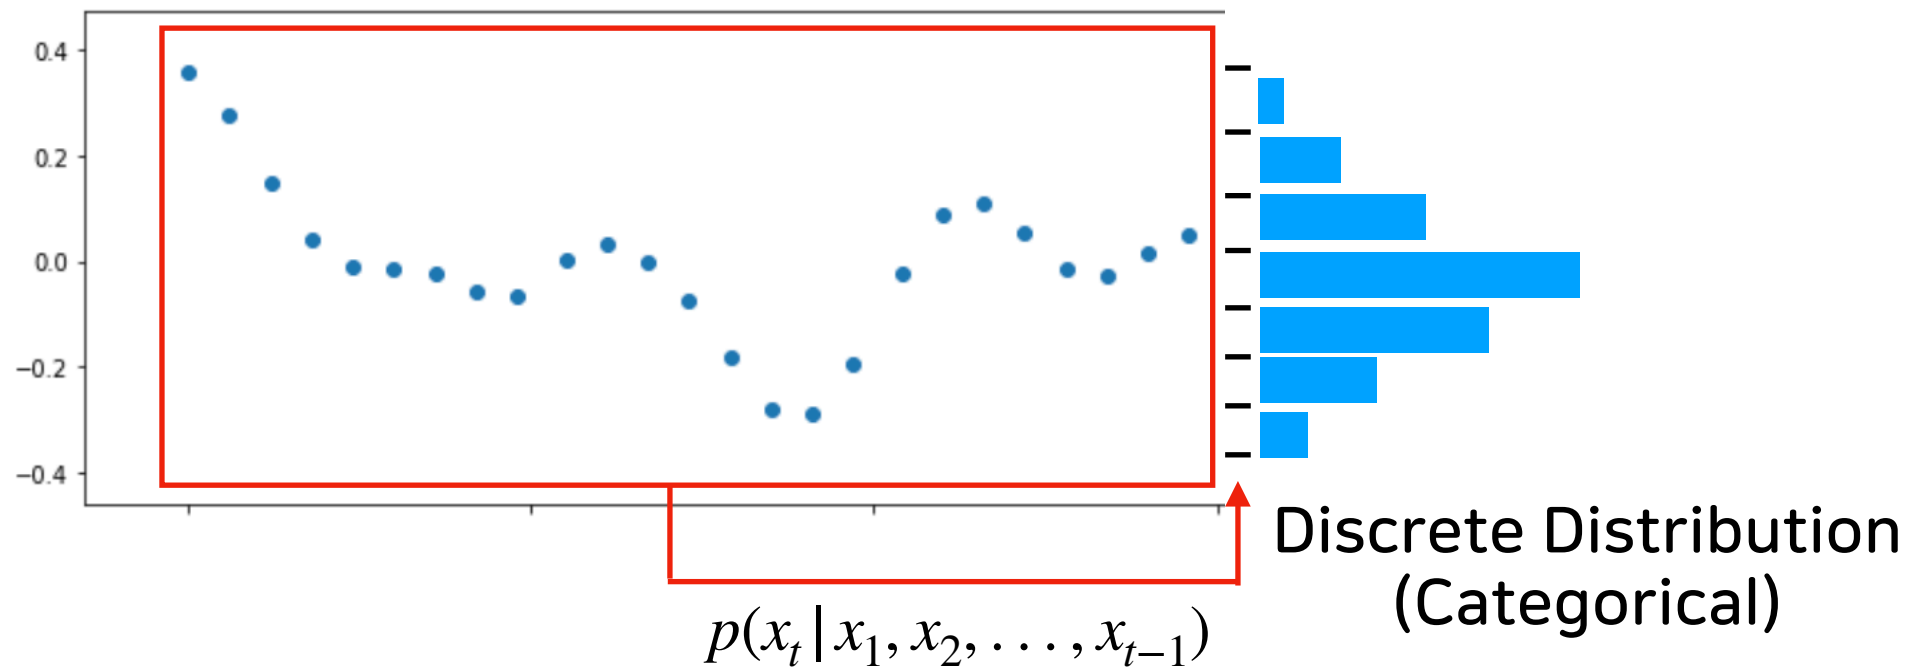
Autoregressive Models

한단계 앞까지의 과거 데이터를 토대로 현재 시점의 데이터의 확률 분포를 예측

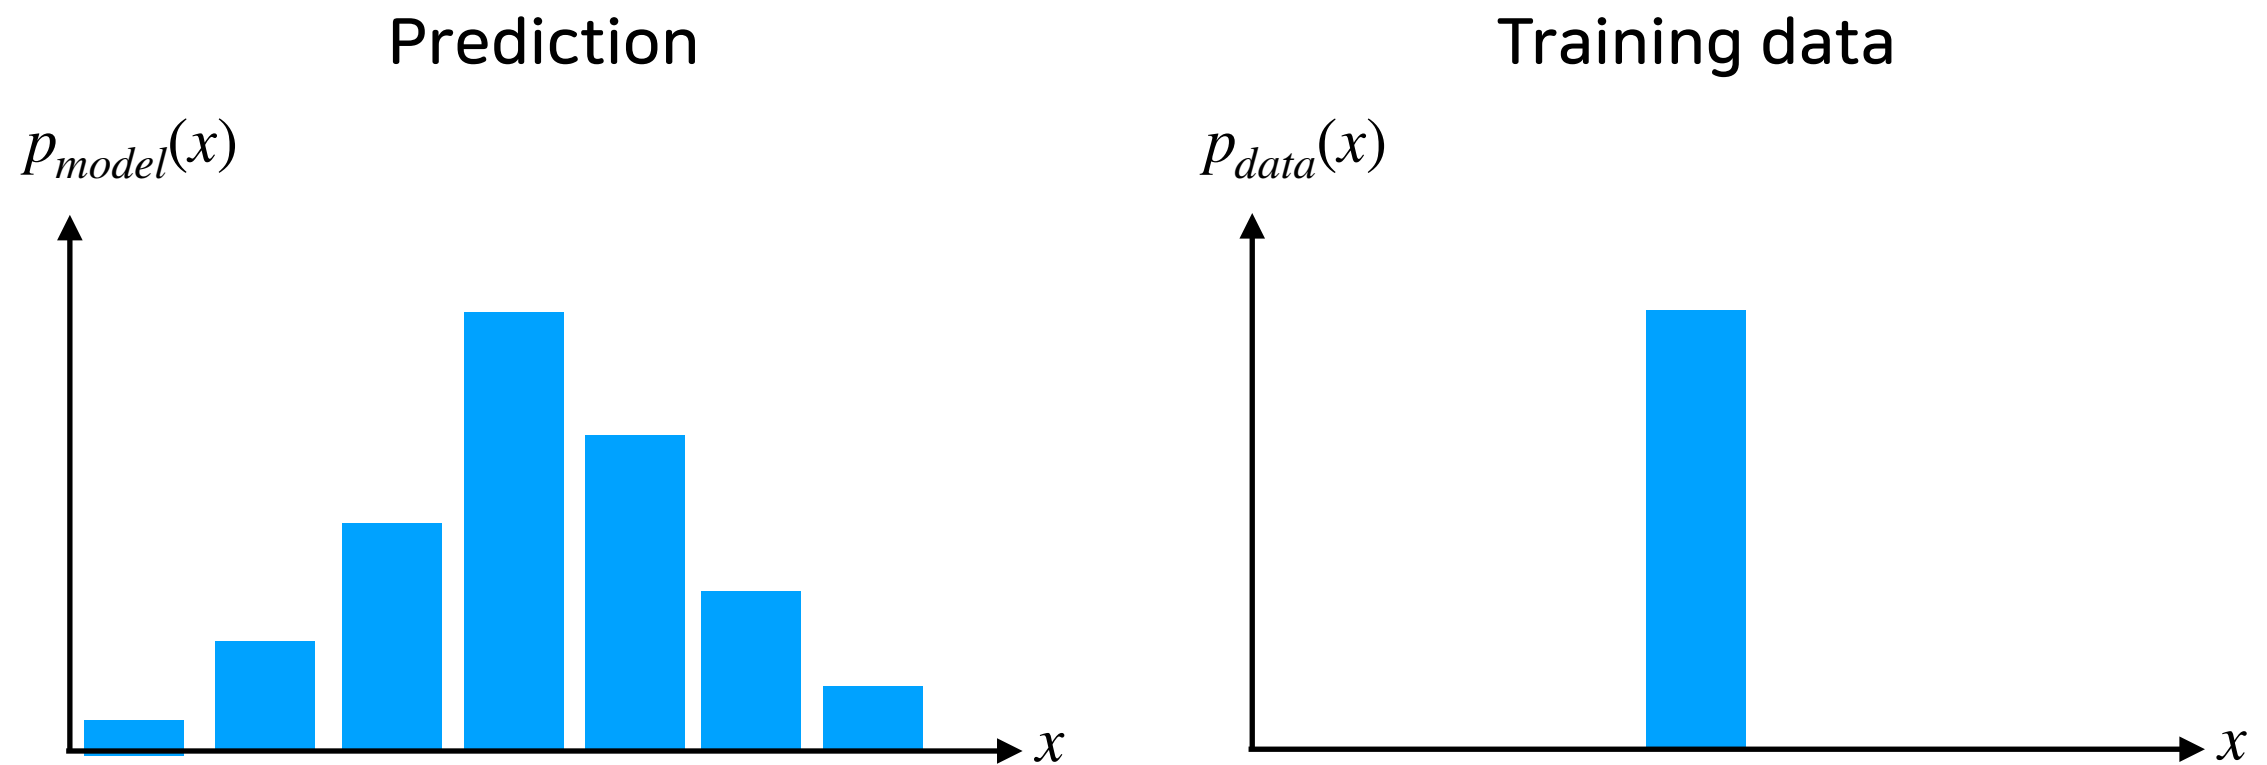


Autoregressive Models

한단계 앞까지의 과거 데이터를 토대로 현재 시점의 데이터의 확률 분포를 예측



Cross-Entropy Loss



$$H(p_{data}, p_{model}) = - \sum_x p_{data}(x) \log p_{model}(x)$$

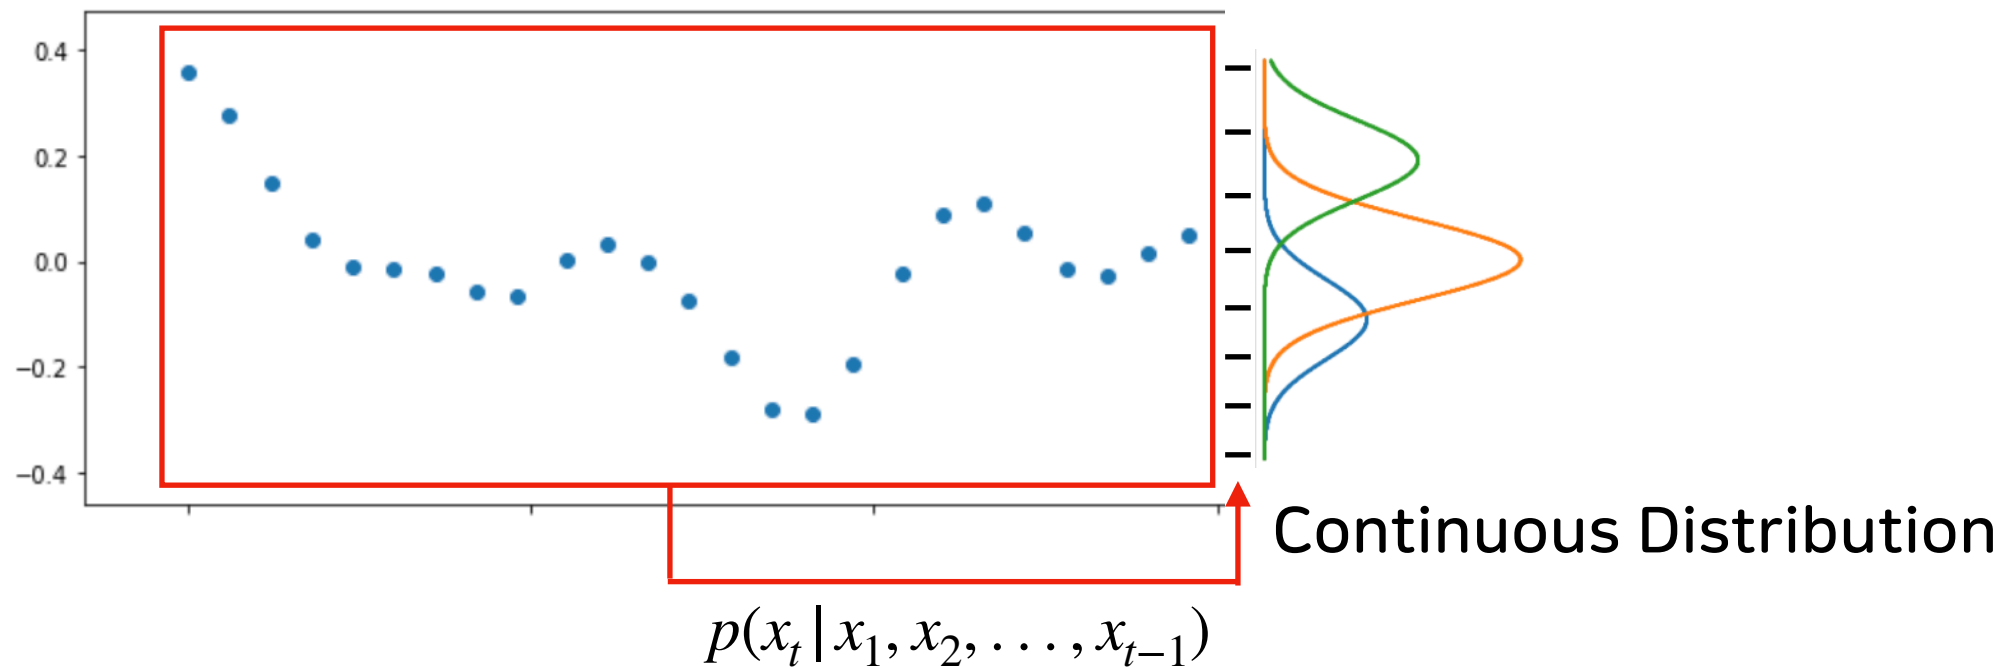
Cross-Entropy Loss

```
346 class MaskedCrossEntropyLoss(nn.Module):
347     def __init__(self):
348         super(MaskedCrossEntropyLoss, self).__init__()
349         self.criterion = nn.CrossEntropyLoss(reduction='none')
350
351     def forward(self, input, target, lengths=None, mask=None, max_len=None):
352         if lengths is None and mask is None:
353             raise RuntimeError("Should provide either lengths or mask")
354
355         # (B, T, 1)
356         if mask is None:
357             mask = sequence_mask(lengths, max_len).unsqueeze(-1)
358
359         # (B, T, D)
360         mask_ = mask.expand_as(target)
361         losses = self.criterion(input, target)
362         return ((losses * mask_).sum()) / mask_.sum()
```

https://github.com/r9y9/wavenet_vocoder/blob/master/train.py

2. Autoregressive Models

한단계 앞까지의 과거 데이터를 토대로 현재 시점의 데이터의 확률 분포를 예측




Chain Rule

$$p(X) = p(x_1, x_2, \dots, x_T) = p(x_1)p(x_2 | x_1)p(x_3 | x_1, x_2) \dots p(x_T | x_1, \dots, x_{T-1})$$

$$= \prod_{t=1}^T p(x_t | x_1, \dots, x_{t-1}) = \prod_{t=1}^T p(x_t | x_{<t})$$

Mixture Density Networks

Figure 5.19 On the left is the data set for a simple ‘forward problem’ in which the red curve shows the result of fitting a two-layer neural network by minimizing the sum-of-squares error function. The corresponding inverse problem, shown on the right, is obtained by exchanging the roles of x and t . Here the same network trained again by minimizing the sum-of-squares error function gives a very poor fit to the data due to the multimodality of the data set.



Mixture Density Networks




Figure 5.20 The *mixture density network* can represent general conditional probability densities $p(\mathbf{t}|\mathbf{x})$ by considering a parametric mixture model for the distribution of \mathbf{t} whose parameters are determined by the outputs of a neural network that takes \mathbf{x} as its input vector.

$$p(\mathbf{t} | \mathbf{x}) = \sum_{k=1}^K \pi_k(\mathbf{x}) \mathcal{N}(\mathbf{t} | \boldsymbol{\mu}_k(\mathbf{x}), \sigma_k^2(\mathbf{x}))$$

Mixture Density Networks

The Mixing coefficients constraints which can be achieved by softmax

$$\sum_{k=1}^K \pi_k(\mathbf{x}) = 1, \quad 0 \leq \pi_k(\mathbf{x}) \leq 1$$

$$\pi_k(\mathbf{x}) = \frac{\exp(a_k^\pi)}{\sum_{l=1}^K \exp(a_l^\pi)}$$

The Variances should satisfy can be represented by the exponentials

$$\sigma_k^2(\mathbf{x}) \geq 0$$

$$\sigma_k(\mathbf{x}) = \exp(a_k^\sigma)$$

The Means have real components

$$\mu_{kj}(\mathbf{x}) = a_{kj}^\mu$$

K : # of components k : component index

Mixture Density Networks


$$Loss(data, \theta) = \mathbb{E}_{p_{data}(\mathbf{x})}[-\log p_{\theta}(\mathbf{x})]$$

$$\approx \sum_n^N [-\log p_{\theta}(\mathbf{x}_n)], \text{ where } \mathbf{x}_n \sim p_{data}(\mathbf{x})$$


$$= - \sum_{n=1}^N \log \left\{ \sum_{k=1}^K \pi_k(\mathbf{x}_n, \theta) \mathcal{N} \left(\mathbf{t}_n | \boldsymbol{\mu}_k(\mathbf{x}_n, \theta), \sigma_k^2(\mathbf{x}_n, \theta) \right) \right\}$$

Mixture Density Networks


Figure 5.21 (a) Plot of the mixing coefficients $\pi_k(x)$ as a function of x for the three kernel functions in a mixture density network trained on the data shown in Figure 5.19. The model has three Gaussian components, and uses a two-layer multi-layer perceptron with five 'tanh' sigmoidal units in the hidden layer, and nine outputs (corresponding to the 3 means and 3 variances of the Gaussian components and the 3 mixing coefficients). At both small and large values of x , where the conditional probability density of the target data is unimodal, only one of the kernels has a high value for its prior probability, while at intermediate values of x , where the conditional density is trimodal, the three mixing coefficients have comparable values. (b) Plots of the means $\mu_k(x)$ using the same colour coding as for the mixing coefficients. (c) Plot of the contours of the corresponding conditional probability density of the target data for the same mixture density network. (d) Plot of the approximate conditional mode, shown by the red points, of the conditional density.




(a)



(b)




(c)



(d)


Logistic Mixture

PDF





$$f(y; \mu, s) = \frac{1}{4s} \operatorname{sech}^2 \left(\frac{x - \mu}{2s} \right)$$

CDF




$$F(y; \mu, s) = \frac{1}{1 + e^{-(y-\mu)/s}}$$





Case.1

$$-0.999 \leq x \leq 0.999$$



$$f(y; \mu, s) \propto F(y + 1/255; \mu, s) - F(y - 1/255; \mu, s) \\ = \sigma((y - \mu + 1/255)/s) - \sigma((y - \mu - 1/255)/s)$$


$$\nu \sim \sum_{i=1}^K \pi_i \text{logistic}(\mu_i, s_i) \quad (1)$$

$$P(x|\pi, \mu, s) = \sum_{i=1}^K \pi_i [\sigma((x + 0.5 - \mu_i)/s_i) - \sigma((x - 0.5 - \mu_i)/s_i)], \quad (2)$$

Tim Salimans et al. PixelCNN++: Improving the PixelCNN with Discretized Logistic Mixture Likelihood and Other Modifications


Case.1

$$-0.999 \leq x \leq 0.999$$



$$f(y; \mu, s) \propto F(y + 1/255; \mu, s) - F(y - 1/255; \mu, s) \\ = \sigma((y - \mu + 1/255)/s) - \sigma((y - \mu - 1/255)/s)$$


Normalization



Case.1

$$-0.999 \leq x \leq 0.999$$


```
58     centered_y = y - means
59     inv_stdv = torch.exp(-log_scales)
60     plus_in = inv_stdv * (centered_y + 1. / (num_classes - 1))
61     cdf_plus = torch.sigmoid(plus_in)
62     min_in = inv_stdv * (centered_y - 1. / (num_classes - 1))
63     cdf_min = torch.sigmoid(min_in)
```



Case.1

$$-0.999 \leq x \leq 0.999$$


```
58     centered_y = y - means
59     inv_stdv = torch.exp(-log_scales)
60     plus_in = inv_stdv * (centered_y + 1. / (num_classes - 1))
61     cdf_plus = torch.sigmoid(plus_in)
62     min_in = inv_stdv * (centered_y - 1. / (num_classes - 1))
63     cdf_min = torch.sigmoid(min_in)
```



Case.1

$$-0.999 \leq x \leq 0.999$$


```
58     centered_y = y - means
59     inv_stdv = torch.exp(-log_scales)
60     plus_in = inv_stdv * (centered_y + 1. / (num_classes - 1))
61     cdf_plus = torch.sigmoid(plus_in)
62     min_in = inv_stdv * (centered_y - 1. / (num_classes - 1))
63     cdf_min = torch.sigmoid(min_in)
```



Case.1


$$-0.999 \leq x \leq 0.999$$

```
65     # log probability for edge case of 0 (before scaling)
66     # equivalent: torch.log(torch.sigmoid(plus_in))
67     log_cdf_plus = plus_in - F.softplus(plus_in)
68
69     # log probability for edge case of 255 (before scaling)
70     # equivalent: (1 - torch.sigmoid(min_in)).log()
71     log_one_minus_cdf_min = -F.softplus(min_in)
72
73     # probability for all other cases
74     cdf_delta = cdf_plus - cdf_min
```



Case.2


$y < -0.999$




$$p(\mathbf{y} < y + 1/255; \mu, s) = F(y + 1/255; \mu, s)$$

Case.2

$$y < -0.999$$




Normalization




Case.2

$$y < -0.999$$

```
65     # log probability for edge case of 0 (before scaling)
66     # equivalent: torch.log(torch.sigmoid(plus_in))
67     log_cdf_plus = plus_in - F.softplus(plus_in)
68
69     # log probability for edge case of 255 (before scaling)
70     # equivalent: (1 - torch.sigmoid(min_in)).log()
71     log_one_minus_cdf_min = -F.softplus(min_in)
72
73     # probability for all other cases
74     cdf_delta = cdf_plus - cdf_min
```




Case.3
 $0.999 < y$




$$p(\mathbf{y} > y - 1/255; \mu, s) = 1 - F(y - 1/255; \mu, s)$$

Case.3
 $0.999 < y$



Normalization

$$p(\mathbf{y} > y - 1/255; \mu, s) = 1 - F(y - 1/255; \mu, s)$$




$$in^- = \frac{y - 1/255 - \mu}{s}$$

$$p(\mathbf{y} > in^-; 0, 1) = F(in^-; 0, 1)$$

Case.3

$0.999 < y$

```
65     # log probability for edge case of 0 (before scaling)
66     # equivalent: torch.log(torch.sigmoid(plus_in))
67     log_cdf_plus = plus_in - F.softplus(plus_in)
68
69     # log probability for edge case of 255 (before scaling)
70     # equivalent: (1 - torch.sigmoid(min_in)).log()
71     log_one_minus_cdf_min = -F.softplus(min_in)
72
73     # probability for all other cases
74     cdf_delta = cdf_plus - cdf_min
```



$$\log(1 - F(in^-; 0, 1)) = \log\left(1 - \frac{1}{1 + e^{-in^-}}\right)$$

$$= \log\left(\frac{1}{1 + e^{in^-}}\right)$$

$$= -\text{softplus}(in^-)$$


$$\log(1 - F(in^-; 0, 1)) = -\text{softplus}(in^-)$$

Case.4

$$\sigma(in^+) - \sigma(in^-) < 1e-5$$

```
76 mid_in = inv_stdv * centered_y
77 # log probability in the center of the bin, to be used in extreme cases
78 # (not actually used in our code)
79 log_pdf_mid = mid_in - log_scales - 2. * F.softplus(mid_in)
```

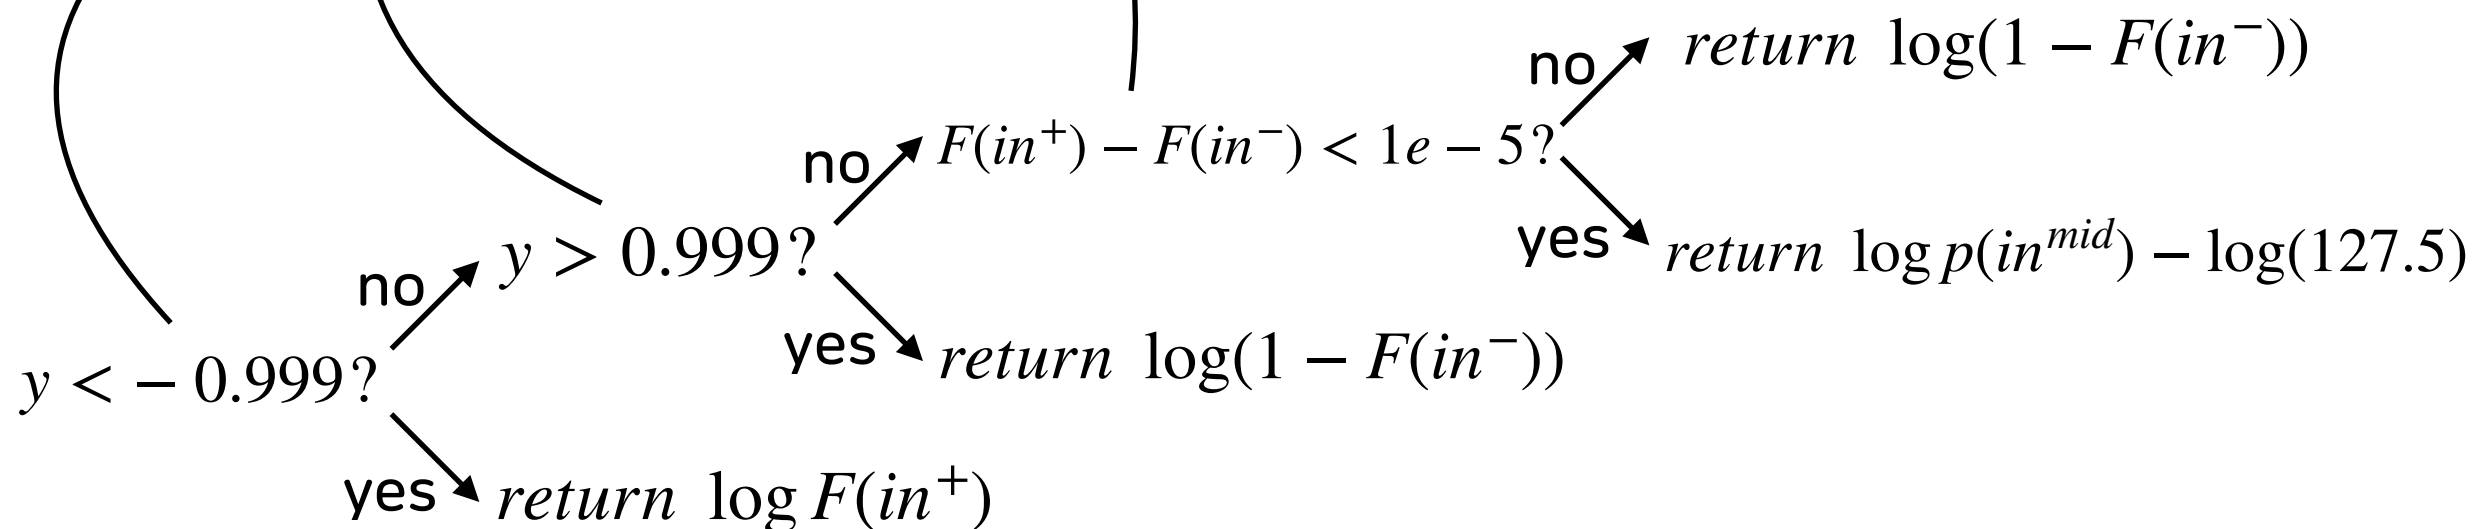
?



```

81 # tf equivalent
82 """
83 log_probs = tf.where(x < -0.999, log_cdf_plus,
84                      tf.where(x > 0.999, log_one_minus_cdf_min,
85                                tf.where(cdf_delta > 1e-5,
86                                          tf.log(tf.maximum(cdf_delta, 1e-12)),
87                                          log_pdf_mid - np.log(127.5))))
88 """
89 # TODO: cdf_delta <= 1e-5 actually can happen. How can we choose the value
90 # for num_classes=65536 case? 1e-7? not sure..
91 inner_inner_cond = (cdf_delta > 1e-5).float()
92
93 inner_inner_out = inner_inner_cond * \
94     torch.log(torch.clamp(cdf_delta, min=1e-12)) + \
95     (1. - inner_inner_cond) * (log_pdf_mid - np.log((num_classes - 1) / 2))
96 inner_cond = (y > 0.999).float()
97 inner_out = inner_cond * log_one_minus_cdf_min + (1. - inner_cond) * inner_inner_out
98 cond = (y < -0.999).float()
99 log_probs = cond * log_cdf_plus + (1. - cond) * inner_out

```



```

101     log_probs = log_probs + F.log_softmax(logits, -1)
102
103     if reduce:
104         return -torch.sum(log_sum_exp(log_probs))
105     else:
106         return -log_sum_exp(log_probs).unsqueeze(-1)

```

$$\begin{aligned}
 p(y) &= \sum_{k=1}^K \pi_k f(y; \mu, s) \\
 NLL = -\log p(y) &= -\log \sum_{k=1}^K \pi_k f(y; \mu, s) \\
 &= -\log \sum_{k=1}^K \exp\{\log \pi_k + \log f(y; \mu, s)\}
 \end{aligned}$$